

	Threat	Best Practice	Deckhouse approach
1	Authentication		
1.1	Unauthenticated access to APIs is provided by the container orchestration tool, allowing unauthorized modification of workloads.	a. All access to orchestration tools components and supporting services – for example, monitoring – from users or other services should be configured to require authentication and individual accountability.	According to the recommended best practice.
1.2	Generic administrator accounts are in place for container orchestration tool management. The use of these accounts would prevent the non-repudiation of individuals with administrator account access.	a. All user credentials used to authenticate to the orchestration should be tied to specific individuals. Generic credentials should not be used. When a default account is present and cannot be deleted, changing the default password to a strong unique password and then disabling the account will prevent a malicious individual from re-enabling the account and gaining access with the default password.	Generic credentials are not used in the platform. Runners are authenticated through service accounts, and users are authenticated through the OIDC authentication mechanism.
1.3	Credentials, such as client certificates, do not provide for revocation. Lost credentials present a risk of unauthorized access to cluster APIs.	a. All credentials used by the orchestration system should be revocable.	Certificates are only used for control plane component authentication in the cluster and never escape control plane nodes.
1.4	Credentials used to access administrative accounts for either containers or container orchestration tools are stored insecurely, leading to unauthorized access to containers or sensitive data.	a. Authentication mechanisms used by the orchestration system should store credentials in a properly secured datastore.	<p>We use Kubernetes secrets or CRs to store authentication credentials. Although those resources are not encrypted by default, the user has all the mechanisms out of the box to do it on their own.</p> <p>For our part, we are thinking about providing a secret store out of the box:</p> <p>https://github.com/deckhouse/deckhouse/issues/3053</p>

1.5	Availability of automatic credentials for any workloads running in the cluster. These credentials are susceptible to abuse, particularly if given excessive rights.	a. Credentials for the orchestration system should only be provided to services running in the cluster where explicitly required.	Zero rights are granted by default. Each application or service gets the least rights required according to the cluster RBAC.
		b. Service accounts should be configured for least privilege. The level of rights they will have is dependent on how the cluster RBAC is configured.	All the granted rights in the cluster are stored in the Deckhouse repo, which eliminates the very possibility of rights injection. There are 6 predefined roles in the cluster with precise rules that we suggest using.
1.6	Static credentials – i.e., passwords – used by administrators or service accounts are susceptible to credential stuffing, phishing, keystroke logging, local discovery, extortion, password spray, and brute force attacks.	a. Interactive users accessing container orchestration APIs should use multi-factor authentication (MFA).	MFA can be connected to the cluster's OIDC, though it is not enabled out of the box. Static passwords aren't used.
2 Authorization			
2.1	Excessive access rights to the container orchestration API could allow users to modify workloads without authorization.	a. Access granted to orchestration systems for users or services should be on a least privilege basis. Blanket administrative access should not be used.	According to the recommended best practice.
2.2	Excessive access rights to the container orchestration tools may be provided through the use of hard-coded access groups.	a. All access granted to the orchestration tool should be capable of modification.	According to the recommended best practice.
		b. Access groups should not be hard-coded.	According to the recommended best practice.
2.3	Accounts may accumulate permissions without documented approvals.	a. Use manual and automated means to regularly audit implemented permissions.	The reconciliation loop is implemented in the platform. All the objects created in the cluster are regularly checked against the desired state.
3 Workload Security			
3.1	Access to shared resources on the underlying host permits container breakouts to occur, compromising the security of shared resources.	a. Workloads running in the orchestration system should be configured to prevent access to the underlying cluster nodes by default.	The Pod Security Standards implemented in the platform define three different policies covering the security spectrum.
		Where granted, any access to resources provided by the	These policies are cumulative and range from highly

		nodes should be provided on a least privilege basis, and the use of "privileged" mode containers should be specifically avoided.	permissive to highly restrictive.
3.2	The use of non-specific versions of container images could facilitate a supply chain attack where a malicious version of the image is pushed to a registry by an attacker.	a. Workload definitions/manifests should target specific known versions of any container images. This should be done via a reliable mechanism checking the cryptographic signatures of images. If signatures are not available, message digests should be used.	We plan to switch to SHA digests for identifying images being deployed: https://github.com/deckhouse/deckhouse/issues/1825
3.3	Containers retrieved from untrusted sources may contain malware or exploitable vulnerabilities.	a. All container images running in the cluster should come from trusted sources.	All container images are stored in the private repo. We build all the open-source platform components on our side from their source code.
4 Network Security			
4.1	Container technologies with container networks that do not support network segmentation or restriction allow unauthorized network access between containers.	a. Container orchestration tool networks should be configured on a default deny basis, with access explicitly required only for the operation of the applications being allowed.	The recommended approach is implemented based on the Cilium network policies. Multitenancy out of the box will be implemented in the platform shortly: https://github.com/deckhouse/deckhouse/issues/1419
4.2	Access from the container or other networks to the orchestration component and administrative APIs could allow privilege escalation attacks.	a. Access to orchestration system components and other administrative APIs should be restricted using an explicit allow-list of IP addresses.	The recommended approach is implemented based on the Cilium network policies. Multitenancy out of the box will be implemented in the platform shortly: https://github.com/deckhouse/deckhouse/issues/1419
4.3	Unencrypted traffic with management APIs is allowed as a default setting, allowing packet sniffing or spoofing attacks.	a. All traffic with orchestration system components APIs should be over encrypted connections, ensuring encryption key rotation meets PCI key and secret requirements.	All the platform components interact over the TLS protocol and are authenticated through certificates.
5 PKI			
5.1	Inability of some container orchestration tool products to support revocation of certificates may lead to misuse of a stolen or lost certificate by attackers.	a. Where revocation of certificates is not supported, certificate-based authentication should not be used.	Certificates are only used for control plane component authentication in the cluster and never escape control plane nodes.
		b. Rotate certificates as required by PCI or customer policies or	

		if any containers are compromised.	
5.2	PKI and Certificate Authority services integrated within container orchestration tools may not provide sufficient security outside of the container orchestration tool environment, which could lead to exploitation of other services that attempt to use this chain of trust.	a. The certificates issued by orchestration tools should not be trusted outside of the container orchestrator environment, as the container orchestrator's Certificate Authority private key can have weaker protection than other enterprise PKI trust chains.	According to the recommended best practice.
6 Secrets Management			
6.1	Inappropriately stored secrets, including credentials, provided through the container orchestration tool, could be leaked to unauthorized users or attackers with some level of access to the environment.	a. All secrets needed for the operation of applications hosted on the orchestration platform should be held in encrypted dedicated secrets management systems.	According to the recommended best practice.
6.2	Secrets stored without version control could lead to an outage if a compromise occurs and there is a requirement to rotate them quickly.	a. Apply version control for secrets, so it is easy to refresh or revoke it in case of a compromise.	According to the recommended best practice.
7 Container Orchestration Tool Auditing			
7.1	Existing inventory management and logging solutions may not suffice due to the ephemeral nature of containers and container orchestration tools integration.	a. Access to the orchestration system API(s) should be audited and monitored for indications of unauthorized access. Audit logs should be securely stored on a centralized system.	The platform collects audit logs; the log-shipper module can then save them to an external storage.
8 Container Monitoring			
8.1	Local logging solutions will not allow for appropriate correlation of security events where containers are regularly destroyed.	a. Centralized logging of container activity should be implemented and allow for correlation of events across instances of the same container.	The log-shipper module implements centralized logging of container activities. The Loki log storage is expected to become part of the platform in the nearest future: https://github.com/deckhouse/deckhouse/issues/2532
8.2	Without appropriate detection facilities, the ephemeral nature of containers may allow attackers to execute attacks unnoticed.	a. Controls should be implemented to detect the adding and execution of new binaries and unauthorized modification of container files to running containers.	The log-shipper logs all cluster activities in the audit logs; Runtime-audit-engine module implements a runtime threats detection engine.
9 Container Runtime Security			

9.1	The default security posture of Linux process-based containers provides a large attack surface using a shared Linux kernel. Without hardening, it may be susceptible to exploits that allow for container escape.	a. Where high-risk workloads are identified, consideration should be given to using either container runtimes that provide hypervisorlevel isolation for the workload or dedicated security sandboxes.	According to the recommended best practice. Open Policy Agent is integrated to check if a pod has privileged or excessive rights.
9.2	Windows process-based containers do not provide a security barrier (per Microsoft's guidance) allowing for possible container break-out.	a. Where Windows containers are used to run application containers, Hyper-V isolation should be deployed in-line with Microsoft's security guidance.	Windows containers are not supported.
10 Patching			
10.1	Outdated container orchestration tool components can be vulnerable to exploits that allow for the compromise of the installed cluster or workloads.	a. All container orchestration tools should be supported and receive regular security patches, either from the core project or backported by the orchestration system vendor.	We regularly check CVEs, renew base images for all the components, and patch Kubernetes to the latest minor versions.
10.2	Vulnerabilities present on container orchestration tool hosts (commonly Linux VMs) will allow for compromise of container orchestration tools and other components.	a. Host operating system of all the nodes that are part of a cluster controlled by a container orchestration tool should be patched and kept up to date. With the ability to reschedule workloads dynamically, each node can be patched one at a time, without a maintenance window.	The platform is not responsible for security updates at the operating system level.
10.3	As container orchestration tools commonly run as containers in the clusters, any container with vulnerabilities may allow compromise of container orchestration tools.	a. All container images used for applications running in the cluster should be regularly scanned for vulnerabilities, patches should be regularly applied, and the patched images redeployed to the cluster.	We use the trivy component to check for vulnerabilities on a daily basis.
11 Resource Management			
11.1	A compromised container could disrupt the operation of applications due to excessive use of shared resources.	a. All workloads running via a container orchestration system should have defined resource limits to reduce the risk of "noisy neighbors" causing availability issues with workloads in the same cluster.	Open Policy Agent controls resource requests and limits for all running containers. Multitenancy will be implemented and will allow you to configure the default resource amounts: https://github.com/deckhouse/deckhouse/issues/1419
12 Container Image Building			
12.1	Container base images downloaded from untrusted sources, or	a. Application container images should be built from trusted,	We build all the platform images on our own.

	which contain unnecessary packages, increase the risk of supply chain attacks.	up-to-date minimal base images.	
12.2	Base images downloaded from external container image registries can introduce malware, backdoors, and vulnerabilities.	a. A set of common base container images should be maintained in a container registry that is under the entity's control.	All platform images are stored in the private registry.
12.3	The default position of Linux containers, which is to run as root, could increase the risk of a container breakout.	a. Container images should be built to run as a standard (non-root) user.	Container images are checked by linter on the platform level.
13 Registry			
13.1	Unauthorized modification of an organization's container images could allow an attacker to place malicious software into the production container environment.	a. Access to container registries managed by the organization should be controlled.	License tokens allow us to control access to the image registry.
		b. Rights to modify or replace images should be limited to authorized individuals.	According to the recommended best practice.
13.2	A lack of segregation between production and non-production container registries may result in insecure images deployed to the production environment.	a. Consider using two registries, one for production or business-critical workloads and one for development/test purposes, to assist in preventing image sprawl and the opportunity for an unmaintained or vulnerable image being accidentally pulled into a production cluster.	We build separate images for development and production environments.
13.3	Vulnerabilities can be present in base images, regardless of the source of the images, via misconfiguration and other methods.	a. If available, registries should regularly scan images and prevent vulnerable images from being deployed to container runtime environments.	We use the trivy component to check images on a regular basis.
13.4	Known good images can be maliciously or inadvertently substituted or modified and deployed to container runtime environments.	a. Registries should be configured to integrate with the image build processes such that only signed images from authorized build pipelines are available for deployment to container runtime environments.	The list of all platform images is stored in the Deckhouse container; its manifest is installed from a trusted source. Content-based hashes are used as tags. We plan to switch to SHA digests for identifying images being deployed in the nearest future: https://github.com/deckhouse/deckhouse/issues/1825
14 Version Management			

14.1	Without proper control and versioning of container orchestration configuration files, it may be possible for an attacker to make an unauthorized modification to an environment's setup.	a. Version control should be used to manage all non-secret configuration files.	According to the recommended best practice.
		b. Related objects should be grouped into a single file.	According to the recommended best practice.
		c. Labels should be used to semantically identify objects.	All the platform-deployed components are explicitly labeled.
15 Configuration Management			
15.1	Container orchestration tools may be misconfigured and introduce security vulnerabilities.	a. All configurations and container images should be tested in a production-like environment prior to deployment.	According to the recommended best practice.
		<div>b. Configuration standards that address all known security vulnerabilities and are consistent with industry-accepted hardening standards and vendor security guidance should be developed for all system components, including container orchestration tools.<div><div>i. Address all known security vulnerabilities.</div><div>ii. Be consistent with industry-accepted system hardening standards or vendor hardening recommendations.</div><div>iii. Be updated as new vulnerability issues are identified.</div></div></div>	We use Trivy to check containers and binaries on a regular basis. On top of that, the code is scanned by linters.
16 Segmentation			
16.1	Unless an orchestration system is specifically designed for secure multitenancy, a shared mixed-security environment may allow attackers to move from a low-security to a high-security environment.	a. Where practical, higher security components should be placed on dedicated clusters. Where this is not possible, care should be taken to ensure complete segregation between workloads of different security levels.	Higher security components can be manually deployed outside the cluster. We plan to implement this feature in the future and make it available out of the box. The multitenancy approach will also be added to the platform: https://github.com/deckhouse/deckhouse/issues/1419
16.2	Placing critical systems on the same nodes as general application containers may allow attackers to disrupt the security of the cluster through the use of shared resources on the container cluster node.	a. Critical systems should run on dedicated nodes in any container orchestration cluster.	We have dedicated nodes for the control plane, system needs (such as monitoring), ingress, and workloads.

16.3	Placing workloads with different security requirements on the same cluster nodes may allow attackers to gain unauthorized access to high security environments via breakout to the underlying node.	a. Split cluster node pools should be enforced such that a cluster user of the low-security applications cannot schedule workloads to the high-security nodes.	The workloads can only run on dedicated worker nodes. Different node groups may be created and configured using taints and node selectors according to the security requirements.
16.4	Modification of shared cluster resources by users with access to individual applications could result in unauthorized access to sensitive shared resources.	a. Workloads and users who manage individual applications running under the orchestration system should not have the rights to modify shared cluster resources, or any resources used by another application.	Out-of-the-box roles are created according to the required rights separation, e.g., 'User' or 'Cluster administrator'.

